Proves

R4 Cheng

November 21, 2025

Loss (Cost) function

> aka. how to reward a machine

L_2:

$$J(\vec{w}) = \frac{1}{2} \sum_{(\vec{x},y)\in D} (f_{\vec{w}}(\vec{x}) - y)^2 + \lambda \|\vec{w}\|_2^2$$

> Euclidean norm: $\|\vec{w}\|_2 = \sqrt{w_0^2 + w_1^2 + \cdots + w_n^2}$

L_1:

$$J(\vec{w}) = \frac{1}{2} \sum_{(\vec{x},y)\in D} (f_{\vec{w}}(\vec{x}) - y)^2 + \lambda \|\vec{w}\|_1^2$$

> Manhattan norm: $\|\vec{w}\|_1 = |w_0| + |w_1| + \cdots + |w_n|$

Target:

$$\min_{\vec{w}} J(\vec{w})$$

(Batch) Gradient Descent

Start from an arbitrary initial weight vector $\vec{w}$.

$$\vec{w'} \leftarrow \vec{w} - \alpha \Delta J(\vec{w}) = \vec{w} - \alpha \frac{\partial}{\partial \vec{w}} J(\vec{w})$$

> α is called learning rate

And for a single training example $(\vec{x}, y)$, we have:

$$\vec{w'} \leftarrow \vec{w} - \alpha(\vec{w} \cdot \vec{x} - y)\vec{x} = \vec{w} + \alpha(y - \vec{w} \cdot \vec{x})\vec{x}$$

Aka. the least mean square (LMS) udpate rule.

> vs perceptron: $\vec{w'} = \vec{w} + y \cdot \vec{x}$

$$
\begin{aligned}
Derivation: \frac{\partial}{\partial \vec{w}} J(\vec{w}) &= \frac{\partial}{\partial \vec{w}} \frac{1}{2} (f(\vec{x}) - y)^2 \\
&= 2 \cdot \frac{1}{2} (f(\vec{x}) - y) \cdot \frac{\partial}{\partial \vec{w}} (f(\vec{x}) - y) \quad // \text{ Chain Rule} \\
&= (\vec{w} \cdot \vec{x} - y) \cdot \frac{\partial}{\partial \vec{w}} (\vec{w} \cdot \vec{x} - y) \\
&= (\vec{w} \cdot \vec{x} - y) \vec{x}
\end{aligned}
$$

Therefore, for the whole dataset D:

$$\vec{w'} \leftarrow \vec{w} + \alpha \sum_{(\vec{x}, y) \in D} (y - \vec{w} \cdot \vec{x}) \vec{x}$$

Stochastic Gradient Descent

Streamline BGD:

$$\vec{w'} \leftarrow \vec{w} + \alpha (y_i - \vec{w} \cdot \vec{x}_i) \vec{x}_i$$

Where $(\vec{x}_i, y_i)$ is a randomly picked training example from D.

Mini-Batch Gradient Descent

> Balance between GD and SGD

$$\vec{w'} \leftarrow \vec{w} + \alpha \sum_{(\vec{x}, y) \in D_m} (y - \vec{w} \cdot \vec{x}) \vec{x}$$

Where $D_m \subset D$ is a randomly picked m examples of training data.

Ridge regression gradient descent

> L_2, becuase it is differentiable

$$\vec{w'} \leftarrow \vec{w} + \alpha \sum_{(\vec{x}, y) \in D} (y - \vec{w} \cdot \vec{x}) \vec{x} - \lambda \vec{w}$$

The function of $\lambda \vec{w}$ is to decrease the magnitude of $\vec{w}$, leading to reduced model complexity and overfitting.

Normal Equation

TODO

Deep Learning

Activation Functions

The activation function introduces non-linearity into the model (network), enabling it to learn and represent complex patterns in the data.

- Sigmoid: $\sigma(x) = \frac{1}{1+e^{-x}} \Rightarrow$ Usually for binary classification questions.

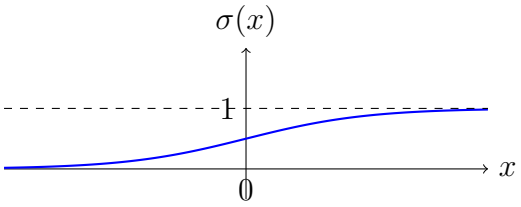

- Hyperbolic Tangent (Tanh): $\tanh(x) = \frac{e^{x}-e^{-x}}{e^{x}+e^{-x}} \Rightarrow$ Make data zero-centered which can speed up convergence.

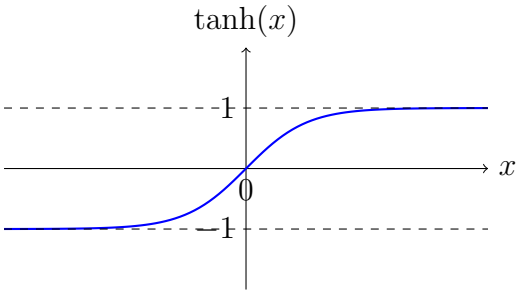

- Rectified Linear Unit (ReLU): $\text{ReLU}(x) = \max(0, x) \Rightarrow$ Only allows positive values to pass through.

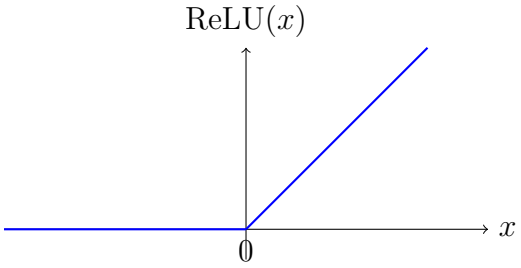

Word Embeddings

Cosine Similarity

Cosine similarity measures the relationship between word vectors in an embedding space.

$$\text{Cosine Similarity}(\mathbf{A}, \mathbf{B}) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\|\|\mathbf{B}\|} = \frac{\sum_{i=1}^{n} A_i B_i}{\sqrt{\sum_{i=1}^{n} A_i^2} \sqrt{\sum_{i=1}^{n} B_i^2}}$$

$$-1 \geq \text{Cosine Similarity}(\mathbf{A}, \mathbf{B}) \leq 1$$

- 1: Vectors are identical (high semantic similarity)
- 0: Vectors are orthogonal (no semantic similarity)
- -1: ?????

Dimensionality Reduction

E.g. 7D to 2D

> Using e.g. PCA, t-SNE, UMAP